


Node in Production

Module 2: Node Production Preparations



Azat Mardan @azat_co



Securing Your App with Environment Variables

Why not store in code?

- Accidentally publishing sensitive information ¹
- GitHub/Bitbucket/GitLab might get and *were* hacked ²
- Code is less flexible to change in a cloud environment

¹ For example, [In major goof, Uber stored sensitive database key on public GitHub page](#) and [How Homakov hacked GitHub and how to protect your application by Peter Nixey](#)
[Raw](#)

² For example, [GitHub Security Update: Reused password attack](#)

Bad practice:

server.js:

```
const apiKey = 'BFED4C44-69DF-49F5'  
const apiSecret = 'AA719195-43D2-4C32-96B9-BBD6CB9400AA'  
const databaseUrl = 'mongodb://admin:EFC429E2-CB84-458F-97BB-CD9BDCCDB1B1@52-53-154-247:27017'  
  
const dbConnection = db.connect(databaseUrl) // connect to a database  
oauth.get(apiKey, apiSecret, accessCode, callback) // fetch protected resource using OAuth
```

Having config .json files (or .js) is still not good since most likely they are stored together with the code

Bad practice:

- config.json
- api/server/datasources.json (Loopback)
- source code (.js)

Use Env Vars

One-liner:

```
MY_ENV_VAR=hello node -e "console.log(`I see: ${process.env.MY_ENV_VAR}`)"
```

In session:

```
export MY_ENV_VAR=hello  
node -e "console.log(`I see: ${process.env.MY_ENV_VAR}`)"
```

Use Env Vars (cont)

server.js:

```
const apiKey = process.env.API_KEY
const apiSecret = process.env.API_SECRET
const databaseUrl = process.env.DATABASE_URL

const dbConnection = db.connect(databaseUrl) // connect to a database
oauth.get(apiKey, apiSecret, accessCode, callback) // fetch protected resource using OAuth
```

A lot of libraries read from NODE_ENV

NODE_ENV=production

NODE_ENV=development

Ways to populate env vars:

- AWS EC2 User Data
- Dockerfile
- AWS CloudFormation blueprint and/or parameters
- CLI or Web interfaces of your cloud provider (AWS, Heroku, etc.)

Separating Development and Production Dependencies

The more dependencies you have, the higher the risk of an attack or a bug. Why have dependencies which you don't need?

Separate dependencies which you only need
for development and don't ship them to
production!

```
"dependencies": {
  "babel-register": "6.11.6",
  "body-parser": "1.13.2",
  "compression": "1.5.1",
  "errorhandler": "1.4.1",
  "express": "4.13.1",
  "express-validator": "2.13.0",
  "hbs": "4.0.0",
  "mongodb": "2.2.25",
  "morgan": "1.6.1",
  "react": "15.5.4",
  "react-dom": "15.5.4"
},
"devDependencies": {
  "axios": "0.13.1",
  "babel-core": "6.10.4",
  "babel-jest": "13.2.2",
  "babel-loader": "6.4.1",
  "babel-preset-react": "6.5.0",
  "node-dev": "3.1.3",
  "webpack": "2.4.1"
}
```

Install dependencies + devDependencies:

```
npm i
```

Install only dependencies:

```
npm i --production
```

Clear devDependencies:

```
npm prune --production
```

Or remove all and reinstall:

```
rm -rf node_modules
```

```
npm i --production
```

Locking in Dependencies Versions

Exact (recommended):

```
"express": "4.15.2"
```

Ranges (bad) could be 4.20, 4.15.4, etc.:

```
"express": "^4.15.2"
```

```
"express": "~4.15.2"
```

```
"express": "~4.x"
```

Any latest (very bad) could be 3.0, 4.15, 5.1:

```
"express": "*"
```

More: <https://docs.npmjs.com/misc/semver>

You are not safe yet!

Dependencies of dependencies often are not locked!

```
npm i
code@1.0.0 /Users/spq201/Code/node-in-production/code
├─┬ express@4.15.2
│   └─┬ accepts@1.3.3
│       └─┬ mime-types@2.1.15
│           └─┬ mime-db@1.27.0
│               └─┬ negotiator@0.6.1
```

Solution to lock them all

- [Shrinkwrap](#)
- Committing, e.g., GitHub, GitLabs, Bitbucket
- Packing: Zip -> S3
- [Yarn](#) (not recommended³)
- [Private registry](#)

³ [Avoid yarn for packages for now and fully enjoy its benefits for application development](#)

npm Shrinkwrap

```
npm shrinkwrap  
wrote npm-shrinkwrap.json
```

npm i will read from the npm-shrinkwrap.json file

Example of npm-shrinkwrap.json

```
...
"prop-types": {
  "version": "15.5.8",
  "from": "prop-types@>=15.5.7 <16.0.0",
  "resolved": "https://registry.npmjs.org/prop-types/-/prop-types-15.5.8.tgz"
},
"react": {
  "version": "15.5.4",
  "from": "react@15.5.4",
  "resolved": "https://registry.npmjs.org/react/-/react-15.5.4.tgz"
},
"setimmediate": {
  "version": "1.0.5",
  "from": "setimmediate@>=1.0.5 <2.0.0",
  "resolved": "https://registry.npmjs.org/setimmediate/-/setimmediate-1.0.5.tgz"
},
"ua-parser-js": {
  "version": "0.7.12",
  "from": "ua-parser-js@>=0.7.9 <0.8.0",
  "resolved": "https://registry.npmjs.org/ua-parser-js/-/ua-parser-js-0.7.12.tgz"
},
"whatwg-fetch": {
  "version": "2.0.3",
  "from": "whatwg-fetch@>=0.10.0",
  "resolved": "https://registry.npmjs.org/whatwg-fetch/-/whatwg-fetch-2.0.3.tgz"
}
...
```

Committing to version control systems

Add

```
git add node_modules
```

Rebuild on CI server:

```
npm rebuild
```

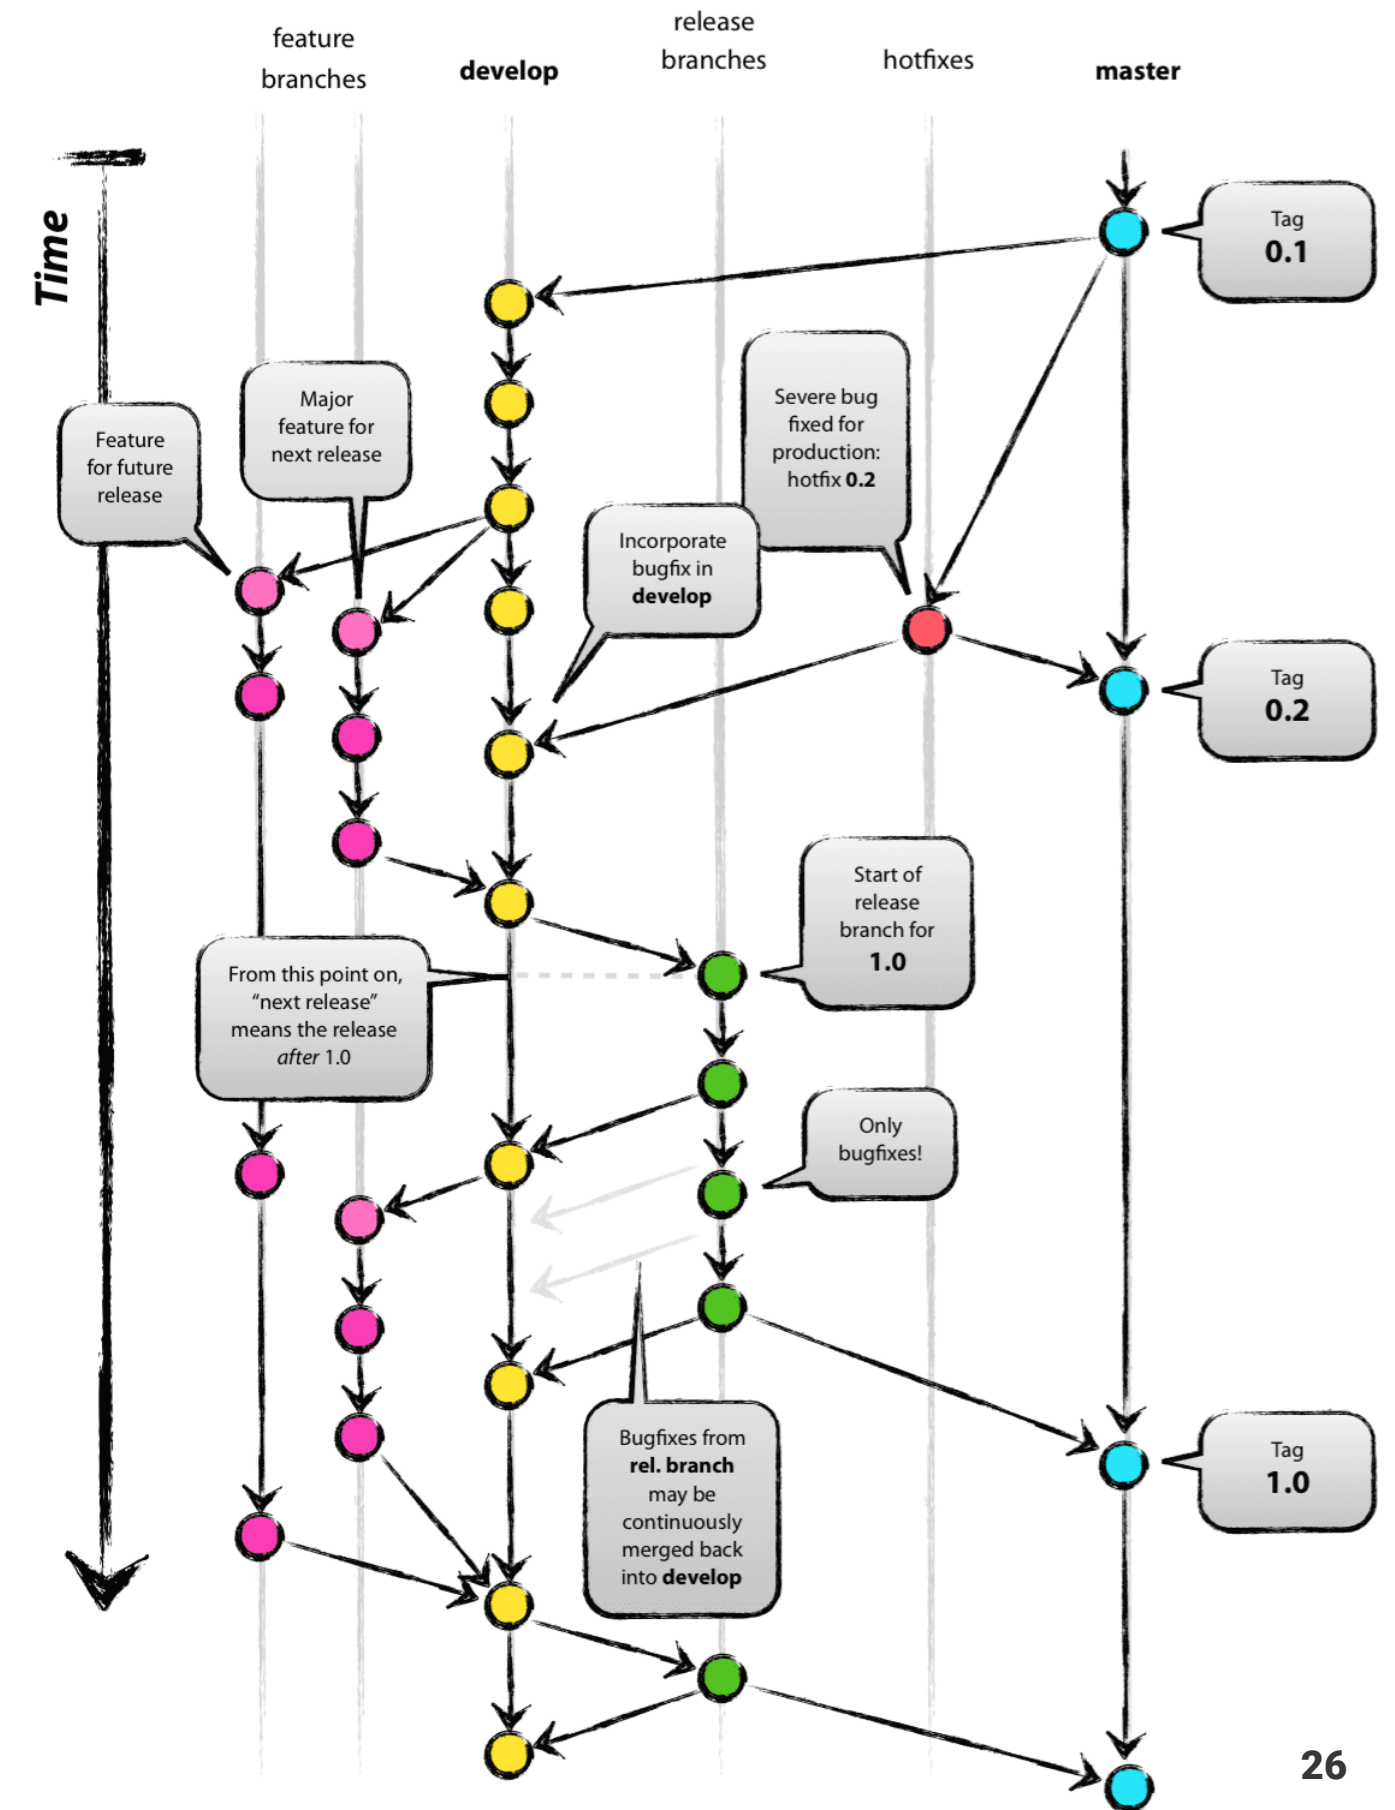
Deployment from Branches

Git flow branches for deployment and releases

- master
- develop
- feature branches
- hotfix branches
- deploy (release)

Source: <http://nvie.com/posts/a-successful-git-branching-model> and <http://nvie.com/files/Git-branching-model.pdf>

Related: <https://hackernoon.com/a-branching-and-releasing-strategy-that-fits-github-flow-be1b6c48eca2> and <https://datasift.github.io/gitflow/IntroducingGitFlow.html>



Stateless Architecture

- Multiple machines in preferably different data centers (AZs)
- Load balancer
- No state in the app layer
- Vertical scaling (pm2, forever, cluster)

Vertical Scaling

- Process managers
- Nginx
- HAProxy
- Varnish Cache
- Redis
- AWS services: ELB, ECS, ElasticCache, Route53, CloudFront, EC2, S3, CloudFormation, AMI

```

var cluster = require('cluster');
var http = require('http');
var numCPUs = require('os').cpus().length;
var express = require('express');
var stats = {};

if (cluster.isMaster) {
  console.log (' Fork %s worker(s) from master', numCPUs)
  for (var i = 0; i < numCPUs; i++) {
    cluster.fork()
  }
  cluster.on('online', function(worker) {
    console.log ('worker is running on %s pid', worker.process.pid)
  })
  cluster.on('exit', function(worker, code, signal) {
    console.log('worker with %s is closed', worker.process.pid)
  })
} else if (cluster.isWorker) {
  var port = 3000
  stats[cluster.worker.process.pid] = 0
  console.log('worker (%s) is now listening to http://localhost:%s',
    cluster.worker.process.pid, port)
  var app = express()
  app.get('*', function(req, res) {
    stats[cluster.worker.process.pid] += 1
    var l = 'cluster '
      + cluster.worker.process.pid
      + ' responded \n'
    console.log(l)
    res.status(200).send(l)
  })
  app.listen(port)
}
process.on('SIGINT', function(){
  console.log(stats)
  console.log('Execute "$ killall node" to terminate')
})

```

Process Managers

- pm2
- StrongLoop PM
- forever - 208 open issues

Install pm2

```
npm i -g pm2@2.4.6
```

Start in development:

```
pm2-dev start server.js
```

Start in production:

```
pm2-docker start -i 0 server.js
```

```
pm2 start -i 0 server.js
```

Consolidated Uogs

- <https://github.com/trentm/node-bunyan>
- <https://github.com/winstonjs/winston>

winston-example.js:

```
const winston = require('winston')

const logger = new (winston.Logger)({
  transports: [
    new (winston.transports.Console)(),
    new (winston.transports.Couchdb)({ 'host': 'localhost', 'db': 'logs' })
    // if you need auth do this:
    // new (winston.transports.Couchdb)({
    //   'user': 'admin', 'pass': 'admin', 'host': 'localhost', 'db': 'logs'
    // })
  ]
})

logger.log('info', 'Hello log files!', { 'message': 'CouchDB logs' })
```

Health-Check Endpoint

GET /status

Response: memory/server load, version SHA, other data.

server.js:

```
app.get('/status', require('routes/status.js'))
```

Node system info

```
const os = require('os')  
os.freemem()  
os.hostname()
```

Node process info

```
process.memoryUsage()  
process.uptime()  
process.version  
process.env.NODE_ENV
```

Node any info

```
const cp = require('child_process')  
cp.exec(posixCommand, (e, result)=>{...})
```

Some POSIX command Node can launch⁴:

```
vmstat -SM -s | grep "used swap" | sed -E "s/[^0-9]*([0-9]{1,8}).*/\1/"  
netstat -an | grep :80 | wc -l  
netstat -an | grep :3000 | wc -l
```

⁴ vmstat: http://linuxcommand.org/man_pages/vmstat8.html

code/status.js:

```
var os = require('os'),
    exec = require('child_process').exec,
    async = require('async'),
    started_at = new Date()

module.exports = function(req, res, next) {
  ...
  res.send({
    status: 'up',
    version: server.get('version'),
    sha: server.et('git sha'),
    started_at: started_at,
    node: {
      version: process.version,
      memoryUsage: Math.round(process.memoryUsage().rss / 1024 / 1024)+"M",
      uptime: process.uptime()
    },
    system: {
      loadavg: os.loadavg(),
      freeMemory: Math.round(os.freemem()/1024/1024)+"M"
    },
    228
    env: process.env.NODE_ENV,
    hostname: os.hostname(),
    connections: connections,
    swap: swap
  })
  ...
}
```


Horizontal Scaling

- AWS ELB: Classic and App ELB
- ECS clusters in 2+ AZs
- S3 for static files
- RDS, DynamoDB or ElasticCache

CI/CD

- AWS CodeDeploy, CodePipeline and build (see [AWS Intermediate](#))
- CircleCI and TravisCI - SaaS
- Drone and Jenkins - hosted

Lab 0: Installs

- Slides, labs and code <https://github.com/azat-co/node-in-production>
- Node and npm (v6 and v4)
- Docker engine
- AWS account and AWS CLI

Detailed instructions and links are in labs/0-installs.md

Time: 15 minutes to download and install, go! 🚀

End of Module 2. Continue with Docker and
AWS 🙋